

AI Deployment Risk Clarity Review

CIEA-powered public sample - fictional case - not legal/compliance certification

HUMAN-GATED / LIMITED GO

For teams that know AI is useful, but cannot yet name what could go wrong.
This review turns unnamed deployment discomfort into a decision-ready risk map: verification, escalation, rollback, accountability, and user/frontline burden.

Fictional case

Acme Support AI

Customer-facing LLM support assistant



Overall readiness
Limited go only

Claimed improvement

Reduce response time and support cost with AI-assisted first response.

Preview finding

Ready to answer faster.
Not ready to own the outcome.

Faster answers do not prove resolution. Hidden work remains in verification, escalation, rollback, and accountable ownership.

1 Who verifies?

Named owner needed.

2 Who escalates?

Clear triggers needed.

3 Who fixes?

Correction path needed.

Readiness dimensions scored

Higher score = more ready / better controlled

Output reliance control 28/100



Low control: customer-facing output invites reliance.

Verification readiness 32/100



Verification work not fully owned or funded.

Rollback readiness 41/100




Partial recovery path; not tested end-to-end.

Accountability recovery 38/100



Named owner exists, but weakly recoverable.

Customer/frontline protection 35/100



Risk of burden moving to users and staff.

Top 3 named problems

Not generic risk categories - operationally named gaps

1

Reliance-Warranty Gap

The interface may invite practical reliance while responsibility remains limited.

2

Output QA Displacement

Generation is faster, but verification work shifts to users, staff, or downstream experts.

3

Escalation Ownership Gap

No clear owner for unresolved, sensitive, or high-stakes outcomes.

Public preview shows the decision memo form - not the full method.

Full scoring basis, threshold logic, internal checklist, counter-surface table, and remediation sequence are intentionally withheld.

LOCKED

SAMPLE

Sample Executive Memo Excerpt

Fictional AI customer-support assistant - redacted preview

REDACTED PREVIEW

Deployment decision preview

HUMAN-GATED / LIMITED GO

Keep deployment human-gated until a verification owner, escalation triggers, and correction path are defined. Limited Go means restricted deployment with human verification and named escalation ownership.

Context

Acme Support AI is proposed as a first-response assistant for customer support. The business claim is lower response time and reduced support cost.

Preview judgment

The current design does not yet show sufficient verification capacity, escalation ownership, or rollback readiness for relied-upon customer-facing output.

What this review clarifies

- **Safe reliance boundary** What AI output is safe to rely on.
- **Verification ownership** Who checks output first.
- **Escalation trigger** When a case must move to a human.
- **Correction path** How harmful or incorrect output is fixed.
- **Failure ownership** Who owns the failure.

Preview classification

Authority expansion	Elevated - review required
Output reliance	Customer-facing
Rollback readiness	Partial / untested
Accountability owner	Named but weakly recoverable
Decision	Human-gated / limited go

Full review includes

- Delegated authority review
- Output reliance and warranty review
- Rollback readiness check
- Accountability exposure mapping
- Counter-surface audit
- Human-gating recommendation

Use this review when AI looks efficient on paper, but verification, escalation, rollback, or accountability remains unclear.

Scope: practical governance readiness review. Not legal advice, compliance certification, penetration test, model benchmark, or safety guarantee.